

Which Damage Do the Images Support? Evidence-Bounded 3D Damage Estimation on Turbine Blades

Anonymous CVPR submission

Paper ID *****

Abstract

Visual inspection of aero-engine blades must decide whether a change seen in an image is physical damage, where it is on the part, and how deep it goes. Existing methods flag anomalies in 2D, output boxes or voxels, or require a measured 3D scan, and none states whether the available views are sufficient to support a claimed shape change. We study evidence-bounded 3D damage estimation: given a healthy CAD model and a few calibrated images, output a displacement field on the blade surface when the images support one, and abstain when camera error, lighting or a surface-colour change explain the images equally well. Our estimator compares a nuisance-only explanation (per-view pose, light gains, an albedo patch) with the same explanation plus a continuous dent, using a differentiable renderer with a neural shader. Dent hypotheses are proposed by a Rao score test computed for every surface location in one backward pass, and the decision threshold on the resulting fit gain is calibrated on healthy and colour-only blades with a split-conformal guarantee. We release BladeDelta, a benchmark of [TODO: N] rendered inspection cases on [TODO: N] public Rotor37 blade geometries with exact damaged meshes, colour-only negatives, three dent families and controlled capture perturbations, with a sealed test split. [TODO: one sentence with headline test numbers]

1. Introduction

Turbine and compressor blades are inspected with cameras, and the inspector's question is physical: is there a dent, where on the blade is it, and how deep is it? Damage of interest is small. On a blade imaged at 512×512 pixels, a dent of 0.07 mm changes a few hundred pixels by two to seven grey levels. A one-degree camera error, a ten-percent lighting change or a stain changes thousands.

Current vision methods answer an easier question. Industrial anomaly detection localises unusual pixels [1, 18], including under pose and illumination change [3, 11]. Reference-based methods compare a photograph with a CAD model and output boxes [2], or lift per-view scores into an anomaly volume [7]. CAD-conditioned reconstruction fits a deformation field to sparse posed views [17, 21], and leaves measurement-quality damaged surfaces to future work [17]. None of these outputs states whether the images actually support a physical change, as opposed to a nuisance explanation.

Our own earlier experiments made this concrete. A render-and-compare fit on public Rotor37 blade geometry produced false dents on 14 of 14 healthy or colour-only blades under a one-degree camera error, and on unfamiliar dents the best image match selected a worse 3D shape in 8 of 12 depth sweeps. Better image agreement is not better geometry, and a plausible reconstruction can choose the wrong physical explanation.

We therefore formulate evidence-bounded 3D damage estimation. The input is a healthy CAD mesh and K calibrated images; the output is a per-vertex displacement field together with a decision, supported or not supported by these views. Our estimator treats this as a nested hypothesis comparison. H_0 explains the images with nuisance only: per-view pose corrections, light gains and a surface albedo patch. H_1 adds a continuous dent. The fit gain $L(H_0) - L(H_1)$ is thresholded at a value calibrated on negatives, giving a controlled false-change rate.

Making this work required separating the roles of the evidence. Pose is fitted from silhouettes, which do not depend on lighting; the dent is compared against nuisance only on interior surface pixels that are not grazing, not near a silhouette and not on the blade root. Dent hypotheses come from a Rao score test evaluated for every surface location and radius with one backward pass through the renderer, and the dent and colour-patch explanations are searched symmetri-

076 cally, so a colour patch fitted first cannot hide a dent.
077 We also report observability: on our development split,
078 70% of dents placed on camera-facing surfaces are less
079 than half visible in the two input views, and the correct
080 output for those is to abstain.

081 Contributions:

- 082 • The task of evidence-bounded 3D damage estimation
083 from sparse views and healthy CAD, with metrics
084 that separate localisation, shape accuracy, false
085 physical change and abstention on unobservable
086 damage.
- 087 • An estimator combining a differentiable renderer
088 with a neural shader, silhouette-only pose, interior-
089 domain nested hypothesis testing, score-test proposals
090 and split-conformal calibration.
- 091 • BladeDelta, a benchmark built on public Rotor37
092 geometry [4] with exact damaged meshes, colour-
093 only negatives, three dent families that differ from
094 the estimator’s own profile, and controlled camera
095 and lighting perturbations, with a protocol and
096 sealed test split fixed before evaluation.

097 2. Related Work

098 Industrial and 3D anomaly detection. Memory-bank
099 methods such as PatchCore [18] localise anomalous pixels;
100 MVTEC 3D-AD [1] and Anomaly-ShapeNet [13]
101 add measured or synthetic 3D. Pose- and illumination-
102 agnostic settings use Gaussian-splat references [11]
103 [TODO: cite PIAD after verifying the primary source],
104 and M²AD [3] benchmarks view–illumination inter-
105 play. Multi-view settings include RAD [5] and
106 SiM3D [7], which lifts per-view scores into an anomaly
107 volume; recent work warns that adding views can leak
108 information in multi-view evaluation [6]. Synthetic
109 anomalies have been argued sufficient for 3D detection
110 [20], and incomplete-modality methods train with
111 3D and test on RGB [19]. These methods output
112 anomaly scores; they do not estimate a physical surface
113 change or state whether the views support one.

114 Reference CAD and blade inspection. Looking 3D [2]
115 aligns a query photograph with a reference shape and
116 localises abnormal regions with boxes. BladeSynth [8]
117 renders procedural nicks, dents, scratches and corrosion
118 on aero-engine blades for 2D detection. Model-free universal
119 detection reconstructs a normal shape by GAN
120 inversion [15]. We instead output a 3D displacement
121 field on a known design.

122 CAD-conditioned reconstruction. CAD-NeRF [21] re-
123 trieves a CAD model and jointly optimises deformation,
124 pose and appearance from few uncalibrated views.

CADsplat [17] fits a shared deformation field to sparse
posed views with a CAD prior and explicitly leaves
measurement-quality deformed surfaces to future work.
Our free-form deformation baseline follows this family.
The difference is the output: a decision on whether
the deformation is supported, with a calibrated false-
change rate.

Inverse rendering and uncertainty. Inverse rendering
separates shape and material from images [16, 22];
we use differentiable rasterisation [12] with a learned
shader for the same purpose on a known geometry.
Radiance-field uncertainty [9, 10] and observation completeness
[14] quantify where a reconstruction is unreliable.
They do not test a physical-change hypothesis against
nuisance explanations, which is the question an
inspector needs answered.

511 3. Method

Setup. We are given a healthy mesh $\mathcal{M} = (V, F)$ with
vertices $V \in \mathbb{R}^{N \times 3}$, K observed images $I_1, \dots, I_K$ with
nominal intrinsics and extrinsics, and nominal lighting.
We seek a displacement field $D \in \mathbb{R}^{N \times 3}$ on the
airfoil and a decision $s \in \{\text{supported}, \text{not supported}\}$.
Nuisance variables η are a per-view pose correction (a
rotation about the part and a camera-frame translation),
three light gains, and an optional albedo patch.

Differentiable renderer. We rasterise $\mathcal{M} + D$ with nvd-
iffrast [12] using smooth vertex normals. A neural
shader maps the per-pixel world normal n , position x
and view direction ω to one radiance component per
light, $R_k(n, x, \omega)$, and an environment term R_0 .
Because light transport is linear in each light’s intensity,
the pixel colour is $T(a \odot (R_0 + \sum_k g_k R_k))$ for albedo
multiplier a , light gains g and a learned monotone tone
curve T , followed by anti-aliasing and the sensor’s pixel
filter. An unstructured shader that took g as an input
learned to ignore it, and lighting changes then produced
spurious dents. The shader is trained once, on
renders of blade geometries disjoint from all evaluation
geometries, with a paired loss on image differences between
lighting settings. The healthy CAD rendered under the
nominal capture provides a per-pixel correction
 $B = I^{\text{ref}} - \hat{I}(\mathcal{M}; \eta_0)$ that absorbs static shader error.
On a held-out blade its mean error is 0.92 grey levels;
image changes caused by a lighting change correlate with
the reference renderer at 0.90–0.98, and changes caused
by a dent at 0.96–0.98 in the 120° view.

Nuisance fit (H_0). Each view’s pose is initialised by
Euclidean ECC alignment of the observed and refer-

ence images, then refined by matching anti-aliased silhouettes to the observed foreground, coarse to fine. Silhouettes do not depend on lighting, which prevents the pose from compensating for a brightness change. With pose fixed, light gains are fitted on the full image.

Evidence domain. All further comparisons use a fixed interior domain Ω : airfoil pixels at least 6 px from any silhouette or self-occlusion edge, excluding the blade root and pixels where $|n \cdot \omega| < 0.35$. Both hypotheses are evaluated on the same Ω , so their losses are comparable.

Dent model (H_1). A dent is parameterised by a sub-vertex centre c (an anchor vertex plus a tangent offset), depth $\delta \geq 0$, radius r , aspect ratio, orientation, profile exponent p and a signed rim coefficient κ :

$$D_i = -\delta \left[(1 - s_i^2)_+^p - \frac{\kappa}{2} \rho(s_i) \right] n_c, \quad (1)$$

where s_i is the anisotropic normalised distance of vertex i from c in the tangent plane, ρ a ring profile and n_c the centre normal, restricted to same-side airfoil vertices. None of the benchmark’s truth families matches this profile exactly.

Score-test proposals. Let $h \in \mathbb{R}^N$ be an inward normal displacement. One backward pass gives $g = -\partial L_\Omega / \partial h$ at $h = 0$. For a kernel $w_{c,r}$ centred at c with radius r , the normalised score $\langle w_{c,r}, g \rangle / \|w_{c,r}\|$ is the Rao score statistic for a small dent at (c, r) against H_0 . We evaluate it for all candidate centres and two radii with a sparse matrix product, and start H_1 fits from the top- k after non-maximum suppression. Proposals are computed from the nuisance-only residual, and each start is fitted both with and without the H_0 albedo patch, so a colour explanation fitted first cannot hide a dent.

Decision. The fit gain is $\Delta = L_\Omega(H_0) - L_\Omega(H_1) \geq 0$, as H_1 nests H_0 . We report D if $\Delta > \tau$ and $D = 0$ otherwise, with τ the split-conformal $(1-\alpha)$ quantile of Δ on calibration negatives (healthy and colour-only blades), which bounds the false-change rate on exchangeable negatives by α . We use a Charbonnier loss ($\epsilon = 2/255$) summed over Ω ; a Student- t noise model and a wider Charbonnier floor did not improve separation on development data.

4. The BladeDelta Benchmark

Geometry. Healthy blades are public Rotor37 compressor geometries [4] (CC BY-SA 4.0), normalised to

Split	Blades	Healthy	Stain	Dent	Dent + stain
Shader training	13	only for the renderer; never evaluated			
Development	12	12	24	36	24
Calibration	12	24	36	36	24
Test (sealed)	30	30	60	90	60

Table 1. BladeDelta case counts. Each case has 7 rendered views; methods use 2 by default.

unit span and closed with a labelled artificial hub collar. Each mesh has 29,989 vertices. Shapes 0–12 were used during earlier development and only to train the shader. Shapes 13–24 form the development split, 25–36 the calibration split and 37–66 the sealed test split; 67–109 are held in reserve.

Damage. Each blade receives one of: no change, a colour-only patch (a soft darkened or rust-tinted albedo change), a dent, or a dent plus a colour patch (half overlapping). Dents follow one of three families — a cubic profile with a raised rim, a Gaussian, and an anisotropic elliptic profile (aspect 0.4–1) — with depth log-uniform in $[0.0005, 0.006]$ blade spans (about 0.04–0.45 mm on a 74 mm Rotor37 blade) and radius in $[0.04, 0.12]$. Centres face at least one of seven cameras and lie at least $\max(0.1, 1.2r)$ from the airfoil boundary.

Capture. Images are rendered with Blender Cycles (32 samples, denoised, AgX tone mapping, 512×512) from seven cameras on a ring around the blade, under one of four conditions: nominal; camera jitter (yaw, pitch and roll each uniform in $\pm 1.5^\circ$ per view); light jitter (each of three area lights $\times [0.85, 1.15]$); or both. Methods receive only the nominal cameras and lighting. The default input is two views (60° and 120°); four- and seven-view settings are ablations.

Protocol. The split, case lists, metrics and hypotheses were hash-frozen before any test image was rendered; protocol changes are logged with dates. For each damaged case we record observability, the depth-weighted fraction of the dent visible in the input views.

CAD6 transfer split. To test transfer beyond Rotor37 we add six industrial blade designs supplied as STEP solids (airfoil spans 24–70 mm), with platforms, fillets and dovetail roots. The airfoil is the largest B-rep face (two faces on one blade); everything else is excluded from evidence like the Rotor37 collar. Each blade is normalised to unit airfoil span in the Rotor37 frame

and receives 76 cases (456 in total): 30% dents centred where both input cameras see the surface, 15% dents hidden from both, 20% healthy, 15% stain only, 10% dent plus stain and 10% healthy with camera error alone ($\pm 0.5^\circ$, $\pm 1^\circ$, $\pm 1.5^\circ$). Dents are sized in millimetres as an inspector would specify them: depth log-uniform in $[0.1, 1.0]$ mm and radius in $[1, 5]$ mm. The case list was hashed before rendering, and no component — shader, threshold or hyperparameter — was refit on this split.

Metrics. Relative local error is the RMS vertex error over the true dent support divided by that of the unchanged CAD (1: no better than CAD; 0: exact). We also report relative whole-airfoil error, depth error, centre error, detection, false physical change on negatives (any vertex moved more than 2×10^{-5}) and abstention on unobservable dents, with 95% bootstrap intervals resampled by blade shape.

5. Experiments

Baselines. Unchanged CAD reports the healthy mesh. Discrete search re-implements our earlier method: image-guided proposals, a fixed cubic profile of radius 0.08 and depths $\{0.0005, 0.001, 0.0015\}$, at nominal pose. Free-form field fits a per-vertex normal displacement with Laplacian smoothness and ℓ_1 sparsity after the same pose and lighting fit. MLP deformation field replaces the per-vertex values with a coordinate MLP (positional encoding, two hidden layers), the continuous deformation field used by CAD-prior reconstruction. These two baselines reimplement the core idea of CADsplat [17] — a smooth deformation field fit to sparse posed views under a CAD prior — within our renderer; they are not the official model, whose code had not been released at the time of writing. 2D detector, lifted is a U-Net trained on (photo, healthy render) pairs of separate exposed blades to segment dents; detections are back-projected onto the mesh through the nominal cameras. All methods share the renderer, cameras and observations, and each is thresholded by the same calibration procedure on its own score.

Main results. [TODO: Table 1: sealed test results]

Risk-coverage.

Ablations.

Why recall is limited. Recall is set by dent depth and by the false-alarm budget, not by a failure to locate (Table 4). Every reported dent is in the right place,




Figure 1. Risk (share of damage claims that are wrong) against coverage (share of cases with a claim) as the threshold varies, sealed test.

Views	Observable dents	AUROC $\uparrow$	Located % $\uparrow$	Median rel. $\downarrow$
2 views	9/49	0.829	56	0.4

Table 2. More views make more dents observable and the evidence clearer.

Variant	AUROC $\uparrow$	Located % $\uparrow$
Full method	0.855	72
No camera (pose) correction	0.679	72
No stain (colour) model	0.753	83
Whole image instead of interior evidence	0.685	61
No score-test proposals	0.420	6
No low-frequency brightness field	0.846	72

Table 3. Ablations on the development subset (18 observable dents, 18 healthy or stained blades).

True depth (mm)	2 views
< 0.1	1/10 (1 loc.)
0.1–0.2	7/14 (7 loc.)
> 0.2	12/13 (12 loc.)
All visible dents	20/37 (20 loc.)

Table 4. Visible test dents reported (located within 0.1 span) by true depth, at the frozen threshold. Below 0.1 mm a dent changes the image by about one grey level; the method declines rather than guess.

and 12/13 visible dents deeper than 0.2mm are reported; most misses are shallower than 0.1 mm, where the change is about one grey level. In 6 of the 17 misses the dent hypothesis converged on the true location, but

its evidence stayed below $\tau = 1.52$. The threshold is held up by appearance changes: stains account for 8 of the 9 false alarms, with stain-only fit gains up to 4.7 (5.6 for one healthy blade under camera error). Lowering τ to 0.54 would reach 70% recall at a 20% false-alarm rate. The misses are therefore the price of the false-alarm guarantee: an inspector receives few spurious damage claims, and dents below about 0.1 mm are reported as unsupported rather than guessed.

Transfer to unseen blade designs. [TODO: CAD6 table] Table ?? applies every method, unchanged, to the six CAD6 blades (– cases). Separation falls from – on the Rotor37 sealed test to – (transfer gap –), with – of visible dents found, – false alarms and – abstention on hidden dents. The threshold was calibrated on Rotor37 only, so the false-alarm rate here measures how far the guarantee survives a change of geometry, not a property the method was fit to. Per blade (Table ??), the method works best on blade – and worst on blade –. Across blades, airfoil span and detection rate have Spearman $\rho = --$ ($p = --$; six blades, so weak evidence): – of visible dents were found on the three smaller blades and – on the three larger ones. Because dents are specified in millimetres, the same dent is relatively deeper on a small blade; detection tracks depth relative to the airfoil span ($\rho = --$) – depth in millimetres ($\rho = --$), rising from – in the shallowest relative-depth tertile to – in the deepest. Recalibrating the threshold on CAD6 negatives, leave-one-blade-out, gives – found at – false alarms: the cost of transferring a threshold across blade designs.

Physical validation (pending real validation).

Pending real validation. A 2× print of CAD6 blade – will be dented with steel balls (target depths 0.3, 0.7 and 1.2 mm), photographed with marker-based poses before and after denting, and measured by structured-light scan (or a 0.01 mm dial gauge). This section is reserved for those results; the rest of the paper does not depend on it.

Qualitative results and failures. [TODO: Figure: observed images, H_0 residual, recovered vs. true dent for successes, a colour/dent ambiguity case and an unobservable dent.]

Real objects. No public dataset pairs real blade photographs with measured damage geometry, so we use SiM3D [7]: six wooden stools scanned with an industrial structured-light sensor before and after damage,

Objects / defects / negatives	6 / 22 / 31
AUROC (defect vs. undamaged)	0.540
Detected (leave-one-object-out)	0/22
Located within 15 mm	0/22
False alarms	1/31

Table 5. Real objects (SiM3D) scanned before and after damage; truth from the measured scans.

with calibrated photographs and expert defect annotations. The healthy scan replaces the CAD model, the healthy photograph from the same scanner position serves as the reference after 2D alignment, and the measured scan difference is the truth; annotated regions with no measured shape change (paint, contamination) count as negatives. On these textured objects our method does not separate physical defects from undamaged areas (Table 5). The residual between sessions is dominated by wood-grain misregistration and specular changes that the shader, trained on 12 photographs per object, does not model. We report this as a failure case: the evidence test is only as reliable as the appearance model behind it.

6. Limitations and Conclusion

Our positive results use rendered images of public compressor geometry with synthetic damage. On real scanned objects (SiM3D) the method did not separate physical defects from undamaged areas; real-blade validation with measured damage remains future work, and an appearance model that holds on real materials is its main prerequisite. The dent model has one dent per blade, and edge damage such as nicks and material loss is out of scope. The conformal guarantee holds for negatives exchangeable with the calibration set; a new camera rig or lighting setup requires recalibration. The method answers whether these views support a physical change; it does not certify airworthiness or replace an inspector.

[TODO: conclusion paragraph after results]

References

- [1] Paul Bergmann, Xin Jin, David Sattlegger, and Carsten Steger. The MVTec 3D-AD Dataset for Unsupervised 3D Anomaly Detection and Localization. arXiv preprint arXiv:2112.09045, 2021. 1, 2
- [2] Ankan Bhunia, Changjian Li, and Hakan Bilen. Looking 3D: Anomaly Detection with 2D-3D Alignment. arXiv preprint arXiv:2406.19393, 2024. 1, 2
- [3] Yunkang Cao, Yuqi Cheng, Xiaohao Xu, Yiheng Zhang, Yihan Sun, Yuxiang Tan, Yuxin Zhang, Xiaonan Huang, and Weiming Shen. Visual Anomaly

- Detection under Complex View-Illumination Interplay: A Large-Scale Benchmark. arXiv preprint arXiv:2505.10996, 2025. 1, 2
- [4] Fabien Casenave, Brian Staber, and Xavier Roynard. MMGP: a Mesh Morphing Gaussian Process-based machine learning method for regression of physical problems under non-parameterized geometrical variability. arXiv preprint arXiv:2305.12871, 2023. 2, 3
- [5] Xinhai Chang, Kaichen Zhou, Taewhan Kim, Jiadong Zhang, Yang Cao, Chufei Peng, Fangneng Zhan, Hao Zhao, Hao Dong, Kai Ming Ting, and Ye Zhu. RAD: A Dataset and Benchmark for Real-Life Anomaly Detection with Robotic Observations. arXiv preprint arXiv:2410.00713, 2024. 2
- [6] Shang-Fu Chen, Kuan-Chuan Peng, Jhih-Ciang Wu, Wen-Huang Cheng, and Kai-Lung Hua. See More, Detect Less? Taming Information Leakage in Multi-View Anomaly Detection. arXiv preprint arXiv:2608.25168, 2026. 2
- [7] Alex Costanzino, Pierluigi Zama Ramirez, Luigi Lella, Matteo Ragaglia, Alessandro Oliva, Giuseppe Lisanti, and Luigi Di Stefano. SiM3D: Single-instance Multiview Multimodal and Multisetup 3D Anomaly Detection Benchmark. arXiv preprint arXiv:2506.21549, 2025. 1, 2, 5
- [8] M. A. Mohammed Eltoum, Ehtesham Iqbal, Yahya Zweiri, Brain Moyo, and Yusra Abdulrahman. BladeSynth: A High-Quality Rendering-Based Synthetic Dataset for Aero Engine Blade Defect Inspection. Scientific Data, 12:1268, 2025. 2
- [9] Chamuditha Jayanga Galappaththige, Thomas Gottwald, Peter Stehr, Edgar Heinert, Niko Sunderhauf, Dimity Miller, and Matthias Rottmann. Predictive Photometric Uncertainty in Gaussian Splatting for Novel View Synthesis. arXiv preprint arXiv:2603.22786, 2026. 2
- [10] Lily Goli, Cody Reading, Silvia Sellán, Alec Jacobson, and Andrea Tagliasacchi. Bayes’ Rays: Uncertainty Quantification for Neural Radiance Fields. arXiv preprint arXiv:2309.03185, 2023. 2
- [11] Mathis Kruse, Marco Rudolph, Dominik Woiwode, and Bodo Rosenhahn. SplatPose & Detect: Pose-Agnostic 3D Anomaly Detection. arXiv preprint arXiv:2404.06832, 2024. 1, 2
- [12] Samuli Laine, Janne Hellsten, Tero Karras, Yeongho Seol, Jaakko Lehtinen, and Timo Aila. Modular Primitives for High-Performance Differentiable Rendering. arXiv preprint arXiv:2011.03277, 2020. 2
- [13] Wenqiao Li, Xiaohao Xu, Yao Gu, Bozhong Zheng, Shenghua Gao, and Yingna Wu. Towards Scalable 3D Anomaly Detection and Localization: A Benchmark via 3D Anomaly Synthesis and A Self-Supervised Learning Network. arXiv preprint arXiv:2311.14897, 2023. 2
- [14] Han Ling, Xian Xu, Yinghui Sun, and Quansen Sun. OCSplats: Observation Completeness Quantification and Label Noise Separation in 3DGS. arXiv preprint arXiv:2508.01239, 2025. 2
- [15] Jiayu Liu, Shancong Mou, Nathan Gaw, and Yinan Wang. Uni-3DAD: GAN-Inversion Aided Universal 3D Anomaly Detection on Model-free Products. arXiv preprint arXiv:2408.16201, 2024. 2
- [16] Yuan Liu, Peng Wang, Cheng Lin, Xiaoxiao Long, Jiepeng Wang, Lingjie Liu, Taku Komura, and Wenping Wang. NeRO: Neural Geometry and BRDF Reconstruction of Reflective Objects from Multiview Images. arXiv preprint arXiv:2305.17398, 2023. 2
- [17] Kristof Overdulse, Lode Jorissen, and Nick Michiels. CADSplat: Sparse-View 3D Gaussian Splatting Aided by CAD Models for Robust, Photorealistic Digital-Twin Reconstruction. arXiv preprint arXiv:2609.18473, 2026. 1, 2, 4
- [18] Karsten Roth, Latha Pemula, Joaquin Zepeda, Bernhard Schölkopf, Thomas Brox, and Peter Gehler. Towards Total Recall in Industrial Anomaly Detection. arXiv preprint arXiv:2106.08265, 2021. 1, 2
- [19] Wenbo Sui, Daniel Lichau, Josselin Lefèvre, and Harold Phelippeau. Incomplete Multimodal Industrial Anomaly Detection via Cross-Modal Distillation. arXiv preprint arXiv:2405.13571, 2024. 2
- [20] Yihan Sun, Yuqi Cheng, Junjie Zu, Yuxiang Tan, Guoyang Xie, Yucheng Wang, Yunkang Cao, and Weiming Shen. Synthesis4AD: Synthetic Anomalies are All You Need for 3D Anomaly Detection. arXiv preprint arXiv:2604.04658, 2026. 2
- [21] Xin Wen, Xuening Zhu, Renjiao Yi, Zhifeng Wang, Chenyang Zhu, and Kai Xu. CAD-NeRF: Learning NeRFs from Uncalibrated Few-view Images by CAD Model Retrieval. arXiv preprint arXiv:2411.02979, 2024. 1, 2
- [22] Kai Zhang, Fujun Luan, Zhengqi Li, and Noah Snavely. IRON: Inverse Rendering by Optimizing Neural SDFs and Materials from Photometric Images. arXiv preprint arXiv:2204.02232, 2022. 2